



# AI6103 Machine Learning Fundamentals

Boyang Li, Albert

School of Computer Science and Engineering

Nanyang Technological University

# Outline of Lecture 3

- Information Theory
- Basics of Machine Learning
- Linear Regression
- Ridge Regression
- Logistic Regression



# Information Theory

- The scientific study of the quantification, storage, and communication of digital information.
- How can we communicate through a noisy channel?
- How can we encode information into binary form efficiently?
- We will only scrape the surface of this vast and rich subject



# Entropy

- The degree of uncertainty or “chaos / surprise / information” in a random variable

- Expectation of negative log probability

$$H(P(X)) = E[-\log X] = -\int P(x) \log P(x) dx$$

- For discrete variables

$$H(P(X)) = -\sum_i P(X = x_i) \log P(X = x_i)$$

- Example: Fair die with probabilities  $\{1/6, 1/6, 1/6, 1/6, 1/6, 1/6\}$

$$H = -\left(\frac{1}{6} \log \frac{1}{6}\right) \times 6 = 0.78$$

- Bias die with probabilities  $\{1/12, 1/12, 1/12, 1/12, 1/3, 1/3\}$

$$H = -\left(\frac{1}{12} \log \frac{1}{12}\right) \times 4 - \left(\frac{1}{3} \log \frac{1}{3}\right) \times 2 = 0.67$$

More uncertainty when the distribution is closer to uniform.



# Entropy: Relation to Event Encoding

- If we observe 26 letters with equal probability, we can use  $\log_2 26 = -\log_2 \frac{1}{26}$  bits to encode each character.
- No fractional bits, so  $\lceil \log_2 26 \rceil = 5$
- A = 00000, B = 00001, C=00010, D=00011, etc.
- To encode three letters, we need 15 bits.



# Entropy: Relation to Event Encoding

- However, if one letter is more common than others, we can design the encoding such that we use fewer bits for more frequent letters.
- A = 001, B=0001, C=10100, D=10011, etc.
- BAD=000100110011 (12 bits)
- Using fewer bits in expectation because less frequent letters have longer encoding.



# Entropy: Relation to Event Encoding

- $-\log_2 P(A)$  is the “information content” of event  $A$ .
- It is the number of bits we need to tell people that this event happened.
- Its expectation is the entropy.

$$H(P) = E[-\log X] = - \sum_i P(X = x_i) \log P(X = x_i)$$



# Cross-Entropy

- For two probability distributions  $P$  and  $Q$ , the cross-entropy is

$$H(Q, P) = E_Q[-\log P(X)] = -\sum_i Q(X = x_i) \log P(X = x_i)$$

- Interpretation: We designed an encoding scheme for the probability distribution  $P$ . However, the actual distribution is  $Q$ . What is the number of bits we need to encode the information?





# Kullback–Leibler divergence (relative entropy)

- A measure for the differences between distributions

$$KL(P||Q) = E_P \left[ \log \frac{P(X)}{Q(X)} \right] = \sum_i P(X = x_i) \log \frac{P(X = x_i)}{Q(X = x_i)}$$

- Continuous distributions with probability density functions  $P$  and  $Q$

$$KL(P||Q) = \int P(x) \log \frac{P(x)}{Q(x)} dx$$



# KL divergence and cross entropy

$$KL(P||Q) = E_P \left[ \log \frac{P}{Q} \right] = \sum_i P(X = x_i) \log \frac{P(X = x_i)}{Q(X = x_i)}$$

$$\begin{aligned} H(Q, P) &= - \sum_i Q(X = x_i) \log P(X = x_i) \quad \text{Cross entropy here} \\ &= - \sum_i Q(X = x_i) \log P(X = x_i) + \sum_i Q(X = x_i) \log Q(X = x_i) - \sum_i Q(X = x_i) \log Q(X = x_i) \\ &= \sum_i Q(X = x_i) \log \frac{Q(X = x_i)}{P(X = x_i)} + H(Q(X)) \\ &= H(Q) + KL(Q || P) \end{aligned}$$

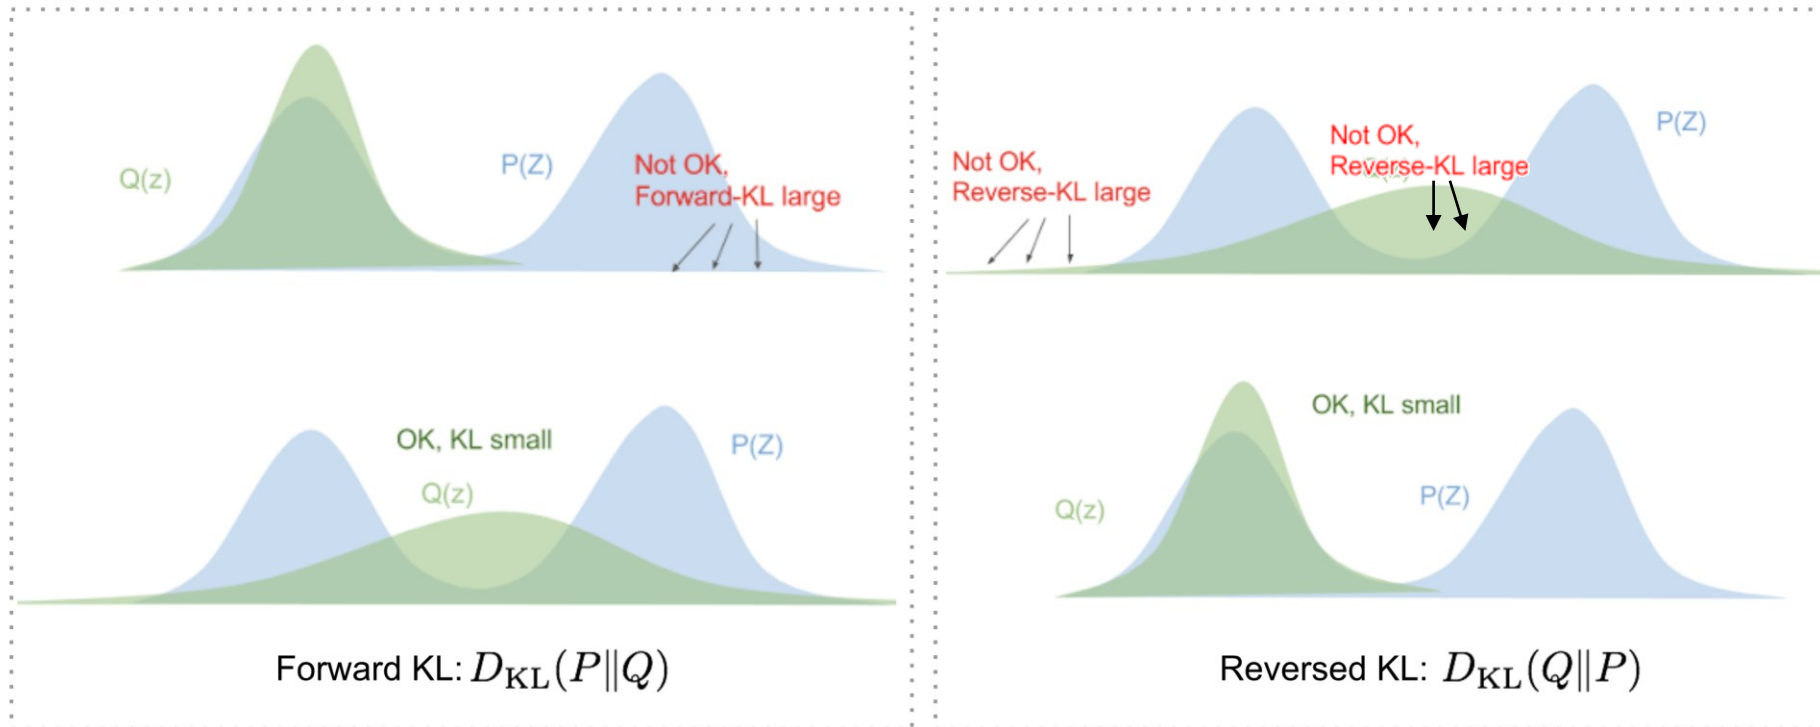


# KL divergence is asymmetric

$$KL(P \parallel Q) = \sum_i P(X = x_i) \log \frac{P(X = x_i)}{Q(X = x_i)}$$

Forward KL:

- P is large when Q is small -> large divergence
- Q is large when P is small -> small divergence.




# Jensen–Shannon divergence

- A symmetric measure for the differences between distributions

$$JS(Q, P) = \frac{1}{2}KL(Q, P) + \frac{1}{2}KL(P, Q)$$



# Machine Learning Basics



# Intelligence?

- Human intelligence has an innate aspect and an environmental aspect
  - Chimpanzees, dolphins, or parrots can demonstrate some levels of intelligence, but they can't reach the human level of intelligence even if training starts very early.
  - As humans, we observe and interact with the world for a long time and learn about knowledge accumulated over thousands of years



# Machine Learning: Analogies

- The “innate” aspect is to specify a machine learning model, which defines the parameters that can be learned and the parameters that are determined before learning (“hyperparameters”).
- The “experiential” aspect is to learn the model parameters from data, a.k.a. training.



# Machine Learning Basics

- There is a function  $y = f(x)$  that we want to approximate
  - $x$  is the input to the machine learning model.
  - $y$  is what the machine learning model tries to predict
- The exact function is unknown, but we have access to historic data  $(\mathbf{x}^{(1)}, y^{(1)}), (\mathbf{x}^{(2)}, y^{(2)}), \dots, (\mathbf{x}^{(n)}, y^{(n)})$
- Our goal is to find out this function from data





# Machine Learning Basics

- There is a function  $y = f(x)$  that we want to approximate
- Example: Automobile insurance risk
  - $x$  = characteristics of the car, such as make, model, year, safety features, engine type, length, weight, height, fuel efficiency, etc.
  - $y$  = probability of accident in a year, or average cost of repairs
- Example: Heart disease diagnosis
  - $x$  = characteristics of the patient, such as age, sex, chest pain location, cholesterol level, blood sugar, etc.
  - $y$  = medical diagnosis made by a human doctor



# Machine Learning Basics

- There is a function  $y = f(x)$  that we want to approximate
- Example: Image classification
  - $x$  = image pixels
  - $y$  = predefined classes, such as dog, cat, truck, airplane, apple, orange, etc.
- Example: Tweet emotion recognition
  - $x$  = text of the tweet
  - $y$  = human label of the reflected emotion: fear, anger, joy, sad, contempt, disgust, and surprise (Ekman's basic emotions).





# Classification vs. Regression

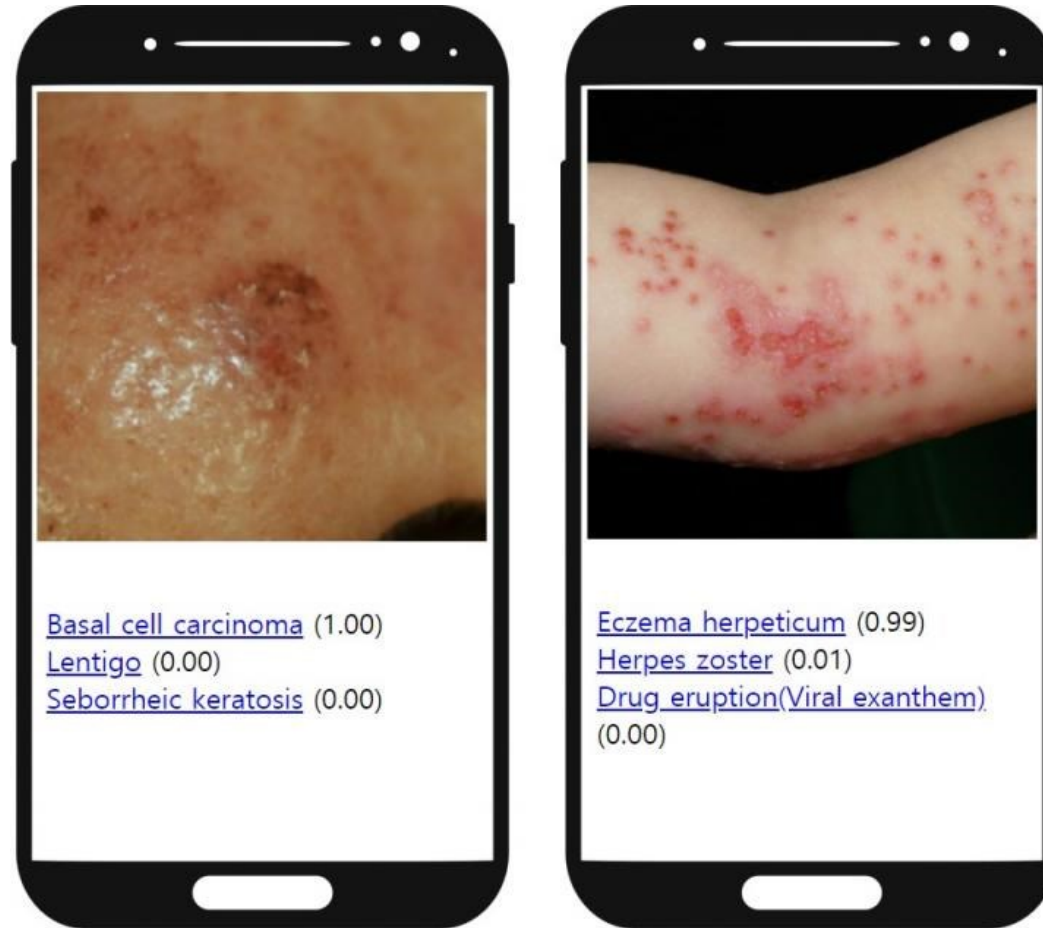
- Classification: the output  $y$  is discrete and represents distinct categories
- Examples
  - Image categories: dog, cat, truck, airplane, apple, orange
  - Emotion categories: fear, anger, joy, sad, contempt, disgust, and surprise



# MNIST Classification



# Classification for Skin Problems




# Classification vs. Regression

- Regression: the output  $y$  represents a continuously varying quantity. Typically, a real number
- Examples
  - Stock price in a week
  - Heart disease: no disease (0), very mild (1), mild (2), severe (3), immediate danger (4)
- Key difference: measurement of error



# Ranking

- Recommender systems provide a list of recommended items.



- We care about not only the first item, but the first N items.



# Useful Features for Wine Quality Rating?

- The following features are probably causal factors to the quality of wine:
  - Alcohol content, pH, residual sugar, free sulfur dioxide, citric acid, tannin, color
- The shape of the bottle is probably unrelated to the quality.
- The year and the winery are probably correlated with the quality.
  - However, we may want to exclude these factors in order to avoid any bias from fame.
- Initial user response may be very good indicators.





# Useful Features for Movie Recommendation?

- The following features are probably useful features for movie recommendation:
  - The list of movies that a user has seen in the past
  - The ratings that the user gave to these movies
  - The user's personal information, such as age, level of education, etc.



# General Guidelines for Feature Selection

- Use features that are strongly correlated with the target variable, but they don't have to be causal.
  - Rooster and sunrise
- Avoid features that you don't want to model to consider, such as the year and the winery in wine quality regression.
  - Beware of information leaks
- Consider data collection, privacy, and ethical concerns



# Linear Regression

- We have a  $p$ -dimensional feature vector  $\mathbf{x} = (x_1, x_2, \dots, x_p)$  and a scalar output  $y$

- Our model is linear

$$\hat{y} = \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_p x_p + \beta_0$$

- In vector form

$$\hat{y} = \boldsymbol{\beta}^\top \mathbf{x} + \beta_0$$



# Linear Regression

- We have  $n$  number of data points

$$(\mathbf{x}^{(1)}, y^{(1)}), (\mathbf{x}^{(2)}, y^{(2)}), \dots, (\mathbf{x}^{(n)}, y^{(n)})$$

- Assumption: Every data point follows the same model

$$\hat{y}^{(i)} = \boldsymbol{\beta}^\top \mathbf{x}^{(i)} + \beta_0$$

- Central question of machine learning

How do we find the parameters  $\boldsymbol{\beta}$  and  $\beta_0$ ?



# One Small Tweak ...

- Adding one dimension to  $\mathbf{x}$ ,

$$\mathbf{x} = (x_1, x_2, \dots, x_p, 1)^\top$$

- Adding one dimension to  $\boldsymbol{\beta}$ ,

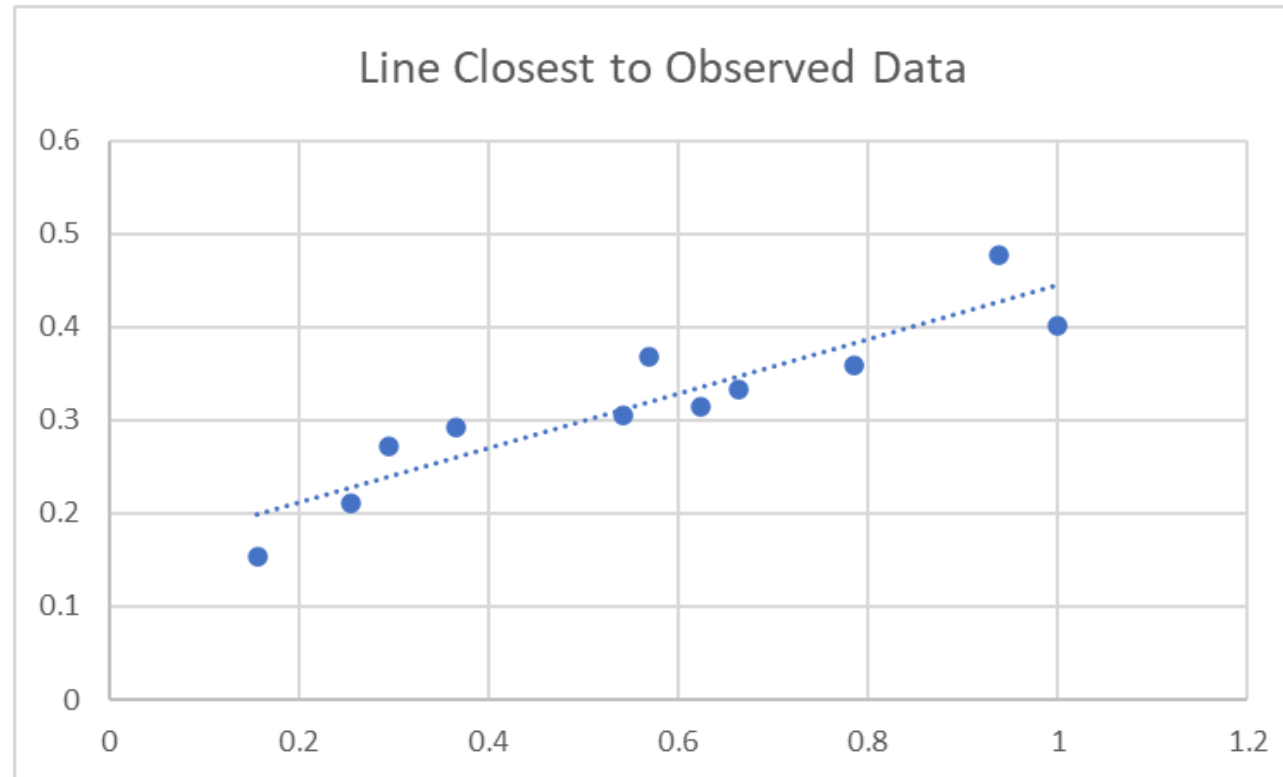
$$\boldsymbol{\beta} = (\beta_1, \beta_2, \dots, \beta_p, \beta_0)^\top$$

- Our model becomes  $\hat{y}^{(i)} = \boldsymbol{\beta}^\top \mathbf{x}^{(i)}$




# Geometric Intuition

- Find the straight line that is closest to observed data points



# Geometric Intuition in 2D

- Find the 2D plane that is closest to observed data points



# Higher dimensions?

- Can't visualize them because we live in a 3D world.
- Geometric intuition: Find the hyperplane that is closest to observed data points
- Key point: The function we fit is linear. A unit change in  $x_i$  always causes a change in  $\hat{y}$  of the magnitude  $\beta_i$ , no matter the value of  $x$  or  $y$ .





# The Loss Function

- We must define a measure of error.
- How wrong is our model?
- Mean Square Error: the average squared distance between  $y^{(i)}$  and  $\hat{y}^{(i)}$

$$\text{MSE} = \frac{1}{n} \sum_{i=1}^n (y^{(i)} - \hat{y}^{(i)})^2 = \frac{1}{n} \|\mathbf{y} - \hat{\mathbf{y}}\|^2$$




# Linear Regression

- One central tenet of supervised learning

**Find the model parameters that lead to smallest error on all possible data.**

- We only observe limited amount of data, so it is usually taken as minimizing error on training data while hoping to achieve low error on test data (data unseen during training)
- When training data are too few or when they are not very representative, we need to use regularization



# A Closed-form Solution

- In matrix form, the loss function is

$$\text{MSE} = \frac{1}{n} (X\beta - y)^\top (X\beta - y)$$

- To find the minimum, we find the derivative against  $\beta$

$$\frac{\partial \text{MSE}}{\partial \beta} = \frac{2}{n} \{X^\top X\beta - X^\top y\}$$

- Necessary condition for minimum: the derivative is zero
- Setting the above to zero and simplify, we get



$$\beta^* = (X^\top X)^{-1} X^\top y$$



# Wait a minute ...

$$\beta = (X^T X)^{-1} X^T Y$$

- Is this a local minimum? What about the second-order condition?



# Second-order Conditions

- Consider the 1D function
$$y = ax^2 + bx + c$$
- When  $a > 0$ , it has a minimum, but no maximum
- When  $a < 0$ , it has a maximum, but no minimum
- When  $a = 0$  and  $b \neq 0$ , it has neither a minimum nor a maximum
- For multidimensional functions, we consider the Hessian matrix.
- This is a symmetric matrix.

$$H = \begin{bmatrix} \frac{\partial^2 L}{\partial \beta_1^2} & \frac{\partial^2 L}{\partial \beta_1 \partial \beta_2} & \cdots & \frac{\partial^2 L}{\partial \beta_1 \partial \beta_p} \\ \frac{\partial^2 L}{\partial \beta_2 \partial \beta_1} & \frac{\partial^2 L}{\partial \beta_2^2} & \cdots & \frac{\partial^2 L}{\partial \beta_2 \partial \beta_p} \\ \vdots & \vdots & \cdots & \vdots \\ \frac{\partial^2 L}{\partial \beta_p \partial \beta_1} & \frac{\partial^2 L}{\partial \beta_p \partial \beta_2} & \cdots & \frac{\partial^2 L}{\partial \beta_p^2} \end{bmatrix}$$



# Linear Algebra: Positive Definiteness

- A square matrix  $A$  is symmetric iff  $A_{ij} = A_{ji}, \forall i, j$

$$\begin{bmatrix} 1 & 3 & 7 \\ 3 & 2 & 6 \\ 7 & 6 & 8 \end{bmatrix} \quad \begin{bmatrix} 1 & 2 & -4 & 38 \\ 2 & -1 & 71 & 2 \\ -4 & 71 & 9 & 56 \\ 38 & 2 & 56 & 30 \end{bmatrix}$$

- A square matrix is anti-symmetric (or skew-symmetric) iff  $A_{ij} = -A_{ji}, \forall i \neq j$

$$\begin{bmatrix} 1 & -3 & -7 \\ 3 & 2 & -6 \\ 7 & 6 & 8 \end{bmatrix}$$

- A matrix  $A$  is positive definite (PD) iff for all vector  $x \neq 0$ ,  $x^T A x > 0$
- Recall  $x^T A x = \sum_i \sum_j A_{ij} x_i x_j$

$$x^T \begin{bmatrix} 2 & -1 & 0 \\ -1 & 3 & -1 \\ 0 & -1 & 2 \end{bmatrix} x = x_1^2 + x_2^2 + x_3^2 + (x_1 - x_2)^2 + (x_3 - x_2)^2 > 0, \forall x \neq 0$$

- Similarly,

$$x^T \begin{bmatrix} 2 & -2 & 0 \\ 0 & 3 & 0 \\ 0 & -2 & 2 \end{bmatrix} x = x_1^2 + x_2^2 + x_3^2 + (x_1 - x_2)^2 + (x_3 - x_2)^2 > 0, \forall x \neq 0$$




# Linear Algebra: Positive Definiteness

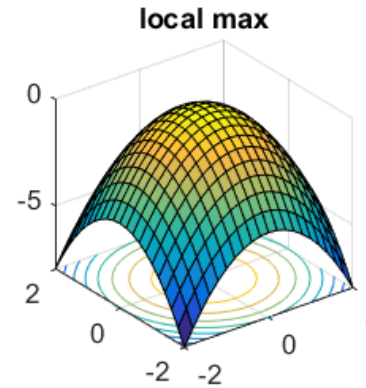
- A matrix  $A$  is positive definite (PD) iff for all vector  $x \neq 0$ ,  $x^\top Ax > 0$
- A matrix  $A$  is positive **semi**-definite (PSD) iff for all vector  $x \neq 0$ ,  $x^\top Ax \geq 0$
- A matrix  $A$  is negative definite (ND) iff for all vector  $x \neq 0$ ,  $x^\top Ax < 0$
- A matrix  $A$  is negative **semi**-definite (NSD) iff for all vector  $x \neq 0$ ,  $x^\top Ax \leq 0$




# Second-order Conditions



Hessian is PSD



Hessian is NSD



Hessian is neither  
PSD nor NSD

$$L = \frac{1}{n} (X\beta - y)^\top (X\beta - y)$$

$$H = \frac{\partial^2 L}{\partial \beta^2} = \frac{2}{n} X^\top X$$





# Conditions for Minimum

- The second-order derivative is  $\frac{\partial^2 \text{MSE}}{\partial \beta^2} = \frac{2}{n} X^\top X$
- It is positive semi-definite because for an arbitrary vector  $z \neq 0$   
$$z^\top X^\top X z = (Xz)^\top Xz$$
- Letting  $a = Xz$ ,  $a^\top a$  is always greater than or equal to zero
- So this is truly a local minimum.
- Since the loss function is convex (which we will not show), the local minimum is also the global minimum.



# Solution to Linear Regression

$$\beta = (X^T X)^{-1} X^T Y$$

is the model parameter that minimizes the loss

$$\begin{aligned} \text{MSE} &= \frac{1}{n} \sum_{i=1}^n (y^{(i)} - \hat{y}^{(i)})^2 \\ &= \frac{1}{n} (X\beta - Y)^T (X\beta - Y) \end{aligned}$$

A.k.a. Ordinary Least Squares



# Another detail: Is $X^T X$ invertible?

$$\beta^* = (X^T X)^{-1} X^T y$$

- When the number of data points  $n$  is greater than the feature dimension  $p$ , and at least  $p$  data points are linearly independent,  $X^T X$  is invertible.
- When we have more features than data points ( $p > n$ ), we have a problem!
- This can be solved by regularization such as ridge regression.



# Ridge Regression

- The ordinary least squares (OLS) estimator:

$$\hat{\beta}_{\text{OLS}} = (X^{\top}X)^{-1}X^{\top}y$$

- If  $n < p$ ,  $X^{\top}X$  is not invertible. We can use ridge regression.

$$\hat{\beta}_{\text{RR}} = (X^{\top}X + \lambda I)^{-1}X^{\top}y$$

- Here  $\lambda$  is a small positive number.
- We can show (but will not) that the ridge regression estimator for  $\beta$  is biased but has lower variance than the OLS estimator [bias-variance tradeoff]



# Ridge Regression Is L2-regularized Linear Regression

- Ridge Regression can be understood as optimizing a different loss function.

$$L = \frac{1}{n} (X\beta - y)^\top (X\beta - y) + \frac{1}{n} \lambda \|\beta\|^2$$

- We again take the derivative against  $\beta$  and set it to zero

$$\frac{\partial L}{\partial \beta} = \frac{2}{n} \{X^\top X\beta - X^\top y + \lambda\beta\} = 0$$

$$\hat{\beta}_{\text{RR}} = (X^\top X + \lambda I)^{-1} X^\top y$$



# Recap: What did we do?

- Collected some data  $(\mathbf{x}^{(1)}, y^{(1)}), (\mathbf{x}^{(2)}, y^{(2)}), \dots, (\mathbf{x}^{(n)}, y^{(n)})$
  - Specified a model  $\hat{y}^{(i)} = \boldsymbol{\beta}^\top \mathbf{x}^{(i)}$
  - Defined a loss function  $\text{MSE} = \frac{1}{n} \sum_{i=1}^n (y^{(i)} - \hat{y}^{(i)})^2$ 
    - Or  $\frac{1}{n} \sum_{i=1}^n (y^{(i)} - \hat{y}^{(i)})^2 + \frac{1}{n} \lambda \boldsymbol{\beta}^\top \boldsymbol{\beta}$
  - Found the parameters  $\boldsymbol{\beta}$  that minimizes the loss function
- Innate, human design
- Experiential, data driven



# Recap: Regularization

- The L2 regularization term  $\beta^\top \beta$  reduces variance of the estimated  $\beta$ .
- This is especially useful when we have limited data.
- We will see many other forms of regularization later.



# Probabilistic Perspective

- The model is parameterized by  $\beta$  and takes input  $x$
- We write its output as  $f_{\beta}(x)$
- We interpret  $f_{\beta}(x)$  as the (input-dependent) parameter  $\mu$  to a Gaussian distribution with unit standard deviation ( $\sigma = 1$ ).

- The ground truth  $y^{(i)}$  is drawn from this distribution

$$y^{(i)} \sim \mathcal{N}(f_{\beta}(x^{(i)}), 1)$$

- Equivalently

$$y^{(i)} = f_{\beta}(x^{(i)}) + \epsilon, \quad \epsilon \sim \mathcal{N}(0, 1)$$





# Maximum Likelihood for Gaussian

- Data:  $y^{(1)}, \dots, y^{(N)}$
- The Gaussian probability is

$$\prod_{i=1}^N P(y^{(i)} | \mu, \sigma) = \prod_{i=1}^N \frac{1}{\sqrt{2\pi}\sigma} \exp - \frac{(y^{(i)} - \mu)^2}{2\sigma^2}$$


- Taking log and remove anything unrelated to  $\mu$

$$\mu^* = \operatorname{argmax}_{\mu} \sum_{i=1}^N \log \frac{1}{\sqrt{2\pi}\sigma} \exp - \frac{(y^{(i)} - \mu)^2}{2\sigma^2} = \operatorname{argmax}_{\mu} \sum_{i=1}^N - \frac{(y^{(i)} - \mu)^2}{2\sigma^2}$$

$$\mu^* = \operatorname{argmin}_{\mu} \sum_{i=1}^N (y^{(i)} - \mu)^2$$



# Plugging in ...



- $\mu^{(i)} = \hat{y}^{(i)} = f_{\beta}(\mathbf{x})$


$$\boldsymbol{\beta}^* = \underset{\boldsymbol{\beta}}{\operatorname{argmin}} \sum_{i=1}^N \left( y^{(i)} - f_{\boldsymbol{\beta}}(\mathbf{x}^{(i)}) \right)^2$$

- Linear regression can be understood as MLE if we assume the label contains noise from the Gaussian distribution.

$$y^{(i)} = f_{\boldsymbol{\beta}}(\mathbf{x}^{(i)}) + \epsilon, \quad \epsilon \sim \mathcal{N}(0, \sigma)$$



# Ridge Regression [Optional]




Ridge regression can be understood as Bayesian maximum a posteriori (MAP) estimation with a Gaussian prior  $\mathcal{N}(0, \frac{1}{\lambda})$  for the model parameters  $\boldsymbol{\beta}$ .


We omit the details for the purpose of this course.





# Linear Models are Limited




# Linear Models are Limited



# Linear Models are Limited



# Non-Linear Models



# Logistic Regression: a Single “Neuron”

- The simplest non-linear model.
- Sometimes referred to as “generalized linear model” as the decision boundary is still linear.
- Here we emphasize the fact that the model function has a non-linear form.





# The Supervised Learning Recipe

- Collect some data  $(\mathbf{x}^{(1)}, y^{(1)}), (\mathbf{x}^{(2)}, y^{(2)}), \dots, (\mathbf{x}^{(n)}, y^{(n)})$
- Specify a model  $\hat{y}^{(i)} = f(\mathbf{x}^{(i)})$
- Define a loss function
- Find the parameters  $\boldsymbol{\beta}$  that minimize the loss function



# The Data

- Collect some data  $(\mathbf{x}^{(1)}, y^{(1)}), (\mathbf{x}^{(2)}, y^{(2)}), \dots, (\mathbf{x}^{(n)}, y^{(n)})$
- $\mathbf{x}^{(i)}$  is a p-vector.
- $y^{(i)}$  is either 0 or 1, denoting the two classes.



# The Supervised Learning Recipe

- Collect some data  $(\mathbf{x}^{(1)}, y^{(1)}), (\mathbf{x}^{(2)}, y^{(2)}), \dots, (\mathbf{x}^{(n)}, y^{(n)})$
- Specify a model  $\hat{y}^{(i)} = f(\mathbf{x}^{(i)})$
- Define a loss function
- Find the parameters  $\beta$  that minimize the loss function



# Logistic Regression: a Single “Neuron”

- Model

$$\hat{y}^{(i)} = \sigma(\boldsymbol{\beta}^\top \mathbf{x}^{(i)})$$

- Activation function

$$\sigma(z) = \frac{1}{1 + e^{-z}} = \frac{e^z}{1 + e^z}$$




# Activation Function: Sigmoid

- Activation function

$$\sigma(z) = \frac{1}{1 + e^{-z}} = \frac{e^z}{1 + e^z}$$

- Squashes all real numbers into the range  $[0, 1]$
- Thus, good for binary classification
- $\sigma(z)$  denotes the probability for one of the two classes.



# Logistic Regression

- Collect some data  $(\mathbf{x}^{(1)}, y^{(1)}), (\mathbf{x}^{(2)}, y^{(2)}), \dots, (\mathbf{x}^{(n)}, y^{(n)})$
- Specify a model  $\hat{y}^{(i)} = \sigma(\boldsymbol{\beta}^\top \mathbf{x}^{(i)})$
- Define a loss function
- Find the parameters  $\boldsymbol{\beta}$  that minimize the loss function



# The Cross-entropy Loss

- The label  $y^{(i)}$  either 0 or 1
- $\hat{y}^{(i)} \in (0, 1)$  is the output of the model.
- The cross-entropy loss

$$L = -\frac{1}{N} \sum_{i=1}^N y^{(i)} \log \hat{y}^{(i)} + (1 - y^{(i)}) \log(1 - \hat{y}^{(i)})$$

For data point  $i$ , only one term exists.



# Why the Name?

- Do they look very similar to you?

$$L_{\text{XE}} = -\frac{1}{N} \sum_{i=1}^N y^{(i)} \log \hat{y}^{(i)} + (1 - y^{(i)}) \log(1 - \hat{y}^{(i)})$$

$$H(P, Q) = E_P[-\log Q(X)]$$





# Why the Name?

- Monte Carlo estimation of an expectation

$$E_P[f(x)] = \int f(x)P(x)dx$$

- The integral can be approximated if we can draw samples  $x^{(1)}, \dots, x^{(K)}$  from  $P(x)$

$$E_P[f(x)] \approx \frac{1}{K} \sum_{i=1}^K f(x^{(i)})$$



# Why the Name?

Cross entropy:

$$H(P, Q) = E_P[-\log Q(X)] = -\sum_i P(X = x_i) \log Q(X = x_i)$$

- $y^{(i)}$  is drawn from an unknown distribution  $P(y^{(i)}|\mathbf{x}^{(i)})$
- $\hat{y}^{(i)}$  is the probability  $Q(y^{(i)} = 1|\mathbf{x}^{(i)}, \boldsymbol{\beta})$
- $1 - \hat{y}^{(i)}$  is the probability  $Q(y^{(i)} = 0|\mathbf{x}^{(i)}, \boldsymbol{\beta})$

$$-\frac{1}{N} \sum_{i=1}^N y^{(i)} \log \hat{y}^{(i)} + (1 - y^{(i)}) \log(1 - \hat{y}^{(i)}) \approx E_P[-\log Q(y^{(i)}|\mathbf{x}^{(i)}, \boldsymbol{\beta})]$$



# Information Theoretical Perspective

- The cross-entropy is related to the KL divergence

$$\begin{aligned} H(P, Q) &= -E_{P(x)}[\log Q(x)] \\ &= H(P) + KL(P||Q) \end{aligned}$$

- Minimizing the loss minimizes the distance between the GT distribution  $P(y^{(i)}|x^{(i)})$  and estimated distribution  $Q(\hat{y}^{(i)}|x^{(i)}, \boldsymbol{\beta})$ .



# MLE Perspective

- Optimizing the cross-entropy can also be seen as the maximum likelihood estimation of  $\beta$  under the binomial distribution.
- We omit the details.



# Logistic Regression

- Collect some data  $(\mathbf{x}^{(1)}, y^{(1)}), (\mathbf{x}^{(2)}, y^{(2)}), \dots, (\mathbf{x}^{(n)}, y^{(n)})$
- Specify a model  $\hat{y}^{(i)} = \sigma(\boldsymbol{\beta}^\top \mathbf{x}^{(i)})$
- Define a loss function: cross entropy
- Find the parameters  $\boldsymbol{\beta}$  that minimize the loss function



# Optimization

- We seek  $\beta$  that minimizes a function  $L(\beta)$
- Assumption: We can evaluate the function and its first-order derivative  $\frac{df(x,w)}{dw}$



# What is a Minimum?

- $x$  is called a local minimum of function  $f(x)$  if there is  $\epsilon$  such that


$$f(x) \leq f(x + y)$$

for all  $\|y\| < \epsilon$ .

- $x$  is called a global minimum of function  $f(x)$  if


$$f(x) \leq f(y)$$

for all  $y$  in the domain of  $f(x)$



# What is a Minimum?


- A global minimum must be a local minimum, but a local minimum may not be a global minimum.
- Multiple local minima cause problems for optimization.






# Optimization: Convex Functions

- The line segment connecting  $f(\mathbf{a})$  and  $f(\mathbf{b})$  always lies above the function between  $\mathbf{a}$  and  $\mathbf{b}$ .
- We can understand a convex function as a function where any local minimum is also a global minimum.
- Easy optimization!



# Optimization: Convex Functions

- Logistic regression has a convex loss function.
- Deep neural networks usually have non-convex loss functions that are difficult to optimize.
- We will worry about that later!




# The Gradient Descent Algorithm

- Input: loss function  $L(\boldsymbol{\beta})$  and initial position  $\boldsymbol{\beta}_0$
- Repeat for a predefined amount of time (or until convergence)
  - Move in the direction of negative gradient
  - $\boldsymbol{\beta}_t = \boldsymbol{\beta}_{t-1} - \eta \left. \frac{dL(\boldsymbol{\beta})}{d\boldsymbol{\beta}} \right|_{\boldsymbol{\beta}_{t-1}}$
- $\eta$  is a small constant called the learning rate



# Gradient Descent

$$\beta_t = \beta_{t-1} - \eta \left. \frac{df(x, \beta)}{d\beta} \right|_{\beta_{t-1}}$$




# Optimization: Gradient Descent

- Starting from a given initial position  $\beta_0$
- Repeat for a predefined amount of time (or until convergence)
  - Move in the direction of negative gradient
  - $\beta_t = \beta_{t-1} - \eta \left. \frac{dL(x, \beta)}{d\beta} \right|_{\beta_{t-1}}$
- This produces a sequence of  $\beta$
- $\beta_0, \beta_1, \dots, \beta_T$
- That goes increasingly closer to the optimum value  $\beta^*$
- If, as  $T \rightarrow \infty$ ,  $\beta_T \rightarrow \beta^*$ , we say that the sequence converges to  $\beta^*$ .





# 2D Functions

- Loss surface in 2D = contour diagrams / level sets  $L_a(f) = \{x | f(x) = a\}$





# All local minima on the diagram



# 2D Functions

- Loss surface in 2D = contour diagrams / level sets  $L_a(f) = \{x | f(x) = a\}$



The gradient direction is the direction along which the function value changes the fastest (for a small change of  $x$  in Euclidean norm).



Along the level set, the function value doesn't change.





# 2D Functions

- Loss surface in 2D = contour diagrams / level sets  $L_a(f) = \{\mathbf{x} | f(\mathbf{x}) = a\}$



For a differentiable function  $f(\mathbf{x})$ , its gradient of  $\frac{df}{d\mathbf{x}}$  at any point is either zero or perpendicular to the level set at that point.




# Gradient Descent on Convex Functions

$$\boldsymbol{\beta}_t = \boldsymbol{\beta}_{t-1} - \eta \left. \frac{df(x, \boldsymbol{\beta})}{d\boldsymbol{\beta}} \right|_{\boldsymbol{\beta}_{t-1}}$$

The learning rate  $\eta$  determines how much we move at each step.

We cannot move too much because the gradient is a local approximation of the function.

Thus, the learning rate is usually small.



Contour diagram / Level sets




# Gradient Descent on Convex Functions

$$\beta_t = \beta_{t-1} - \eta \left. \frac{df(x, \beta)}{d\beta} \right|_{\beta_{t-1}}$$

The learning rate  $\eta$  determines how much we move at each step.

Too small a learning rate  $\eta$ : **slow convergence**

Too large a learning rate  $\eta$ : **oscillation, overshooting**



Contour diagram / Level sets




# Gradient Descent on Convex Functions

$$\beta_t = \beta_{t-1} - \eta \left. \frac{df(x, \beta)}{d\beta} \right|_{\beta_{t-1}}$$

The learning rate  $\eta$  determines how much we move at each step.

As we move closer to the minimum, we often decrease  $\eta$  so that we do not overshoot.



Contour diagram / Level sets




# Logistic Regression

- Collect some data  $(\mathbf{x}^{(1)}, y^{(1)}), (\mathbf{x}^{(2)}, y^{(2)}), \dots, (\mathbf{x}^{(n)}, y^{(n)})$
- Specify a model  $\hat{y}^{(i)} = \sigma(\boldsymbol{\beta}^\top \mathbf{x}^{(i)})$
- Define a loss function: cross entropy
- Find the parameters  $\boldsymbol{\beta}$  that minimizes the loss function



# A single neuron is still **VERY** limited

- It only works well when there is a straight line that can separate two classes.




# A single neuron is still VERY limited


- Perceptron (similar to logistic regression) infamously fails to represent the XOR function (Minsky & Papert, 1969).



$(0, 1)$



$(1, 1)$



$(0, 0)$



$(1, 0)$

